

# 第一次作业

## 查看帮助命令

```
Last login: Fri Mar 24 18:33:01 2023 from 189.80.43.11
[root@masterNode1 ~]# hdfs dfs -help chgrp
-chgrp [-R] GROUP PATH... :
    This is equivalent to -chown ... :GROUP ...
[root@masterNode1 ~]# hdfs dfs -mkdir /file
[root@masterNode1 ~]# hdfs dfs -ls /
Found 1 items
drwxr-xr-x  - root supergroup          0 2023-03-24 18:34 /file
[root@masterNode1 ~]# hdfs dfs -chgrp -R root /file
[root@masterNode1 ~]# hdfs dfs -ls /
Found 1 items
drwxr-xr-x  - root root                0 2023-03-24 18:34 /file
[root@masterNode1 ~]#
```

## 查看新目录所属用户组、修改文件所属用户组

```
[root@masterNode1 ~]# hdfs dfs -mkdir /chmodFile
h[root@masterNode1 ~]# hdfs dfs -ls -R /
drwxr-xr-x  - root supergroup          0 2023-03-24 18:37 /chmodFile
drwxr-xr-x  - root root                0 2023-03-24 18:34 /file
[root@masterNode1 ~]# hdfs dfs -chmod 777 /chmodFile
h[root@masterNode1 ~]# hdfs dfs -ls -R /
drwxrwxrwx  - root supergroup          0 2023-03-24 18:37 /chmodFile
drwxr-xr-x  - root root                0 2023-03-24 18:34 /file
[root@masterNode1 ~]# hdfs dfs -mkdir /file1
# [root@masterNode1 ~]# #查看新目录所属用户组
[root@masterNode1 ~]# hdfs dfs -ls /
Found 3 items
drwxrwxrwx  - root supergroup          0 2023-03-24 18:37 /chmodFile
drwxr-xr-x  - root root                0 2023-03-24 18:34 /file
drwxr-xr-x  - root supergroup          0 2023-03-24 18:38 /file1
[root@masterNode1 ~]# #修改文件所属用户组
[root@masterNode1 ~]# hdfs dfs -chown -R root:SUPERGROUP /file1
[root@masterNode1 ~]# #查看更改文件所属用户或组后的信息
[root@masterNode1 ~]# hdfs dfs -ls /
Found 3 items
drwxrwxrwx  - root supergroup          0 2023-03-24 18:37 /chmodFile
drwxr-xr-x  - root root                0 2023-03-24 18:34 /file
drwxr-xr-x  - root SUPERGROUP          0 2023-03-24 18:38 /file1
```

## 创建一个 datafile 文件夹、在 datafile 里创建一个文件、查看 datafile 中文件的副本数、修改副本数命令

```
[root@masterNode1 ~]# #创建一个datafile文件夹
[root@masterNode1 ~]# hdfs dfs -mkdir /datafile
[root@masterNode1 ~]# #在datafile里创建一个文件
[root@masterNode1 ~]# hdfs dfs -touchz /datafile/Test1.txt
touchz: `/datafile/Test1.txt': No such file or directory
[root@masterNode1 ~]# #查看datafile中文件的副本数
[root@masterNode1 ~]# hdfs dfs -ls -R /datafile
[root@masterNode1 ~]# hdfs dfs -touchz /datafile/Test1.txt
[root@masterNode1 ~]# hdfs dfs -ls -R /datafile
-rw-r--r--    2 root supergroup          0 2023-03-24 18:46 /datafile/Test1.txt
[root@masterNode1 ~]# #修改副本数命令
[root@masterNode1 ~]# hdfs dfs -setrep -w 3 /datafile
Replication 3 set: /datafile/Test1.txt
Waiting for /datafile/Test1.txt ... done
[root@masterNode1 ~]# #查看datafile中文件的副本数
[root@masterNode1 ~]# hdfs dfs -ls -R /datafile
-rw-r--r--    3 root supergroup          0 2023-03-24 18:46 /datafile/Test1.txt
```

## 创建 trashFile.txt 文件、查看 Trash 文件、清空回收站

```
[root@masterNode1 ~]# vim /usr/hadoop-2.7.1/etc/hadoop/core-site.xml
[root@masterNode1 ~]# #创建 trashFile.txt 文件
[root@masterNode1 ~]# hdfs dfs -touchz /datafile/trashFile.txt
[root@masterNode1 ~]# #删除 trashFile.txt 文件
[root@masterNode1 ~]#
[root@masterNode1 ~]# hdfs dfs -rm /datafile/trashFile.txt
23/03/24 18:57:59 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 1440 minutes, Emptier interval = 0 minutes.
Moved: 'hdfs://masterNode1:9000/datafile/trashFile.txt' to trash at: hdfs://masterNode1:9000/user/root/.Trash/Current
[root@masterNode1 ~]#
[root@masterNode1 ~]# #查看 .Trash 文件夹
[root@masterNode1 ~]# hdfs dfs -ls -R /
drwxrwxrwx - root supergroup 0 2023-03-24 18:37 /chmodFile
drwxr-xr-x - root supergroup 0 2023-03-24 18:57 /datafile
-rw-r--r-- 3 root supergroup 0 2023-03-24 18:46 /datafile/Test1.txt
drwxr-xr-x - root root 0 2023-03-24 18:34 /file
drwxr-xr-x - root SUPERGROUP 0 2023-03-24 18:38 /file1
drwx----- - root supergroup 0 2023-03-24 18:57 /user
drwx----- - root supergroup 0 2023-03-24 18:57 /user/root
drwx----- - root supergroup 0 2023-03-24 18:57 /user/root/.Trash
drwx----- - root supergroup 0 2023-03-24 18:57 /user/root/.Trash/Current
drwx----- - root supergroup 0 2023-03-24 18:57 /user/root/.Trash/Current/datafile
-rw-r--r-- 2 root supergroup 0 2023-03-24 18:57 /user/root/.Trash/Current/datafile/trashFile.txt
[root@masterNode1 ~]# #清空回收站
[root@masterNode1 ~]# hdfs dfs -expunge
23/03/24 18:59:19 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 1440 minutes, Emptier interval = 0 minutes.
23/03/24 18:59:19 INFO fs.TrashPolicyDefault: Created trash checkpoint: /user/root/.Trash/230324185919
[root@masterNode1 ~]# #查看 .Trash 文件夹
[root@masterNode1 ~]# hdfs dfs -ls -R /user/root/.Trash/
drwx----- - root supergroup 0 2023-03-24 18:57 /user/root/.Trash/230324185919
drwx----- - root supergroup 0 2023-03-24 18:57 /user/root/.Trash/230324185919/datafile
-rw-r--r-- 2 root supergroup 0 2023-03-24 18:57 /user/root/.Trash/230324185919/datafile/trashFile.txt
```

## 关闭所有服务

```
[root@masterNode1 ~]# #关闭所有服务
[root@masterNode1 ~]# stop-all.sh
This script is Deprecated. Instead use stop-dfs.sh and stop-yarn.sh
Stopping namenodes on [masterNode1]
masterNode1: Warning: Permanently added 'masterNode1,100.100.104.48' (ECDSA) to the list of known hosts.
masterNode1: stopping namenode
masterNode1: Warning: Permanently added 'masterNode1,100.100.104.48' (ECDSA) to the list of known hosts.
slaveNode2: Warning: Permanently added 'slaveNode2,100.100.105.20' (ECDSA) to the list of known hosts.
slaveNode1: Warning: Permanently added 'slaveNode1,100.100.106.13' (ECDSA) to the list of known hosts.
masterNode1: stopping datanode
slaveNode2: stopping datanode
slaveNode1: stopping datanode
Stopping secondary namenodes [0.0.0.0]
0.0.0.0: Warning: Permanently added '0.0.0.0' (ECDSA) to the list of known hosts.
0.0.0.0: stopping secondarynamenode
stopping yarn daemons
stopping resourceManager
masterNode1: Warning: Permanently added 'masterNode1,100.100.104.48' (ECDSA) to the list of known hosts.
slaveNode2: Warning: Permanently added 'slaveNode2,100.100.105.20' (ECDSA) to the list of known hosts.
slaveNode1: Warning: Permanently added 'slaveNode1,100.100.106.13' (ECDSA) to the list of known hosts.
masterNode1: stopping nodemanager
slaveNode2: stopping nodemanager
slaveNode1: stopping nodemanager
no proxyserver to stop
```

## 开启 hadoop

```
[root@masterNode1 ~]# #修改hadoop配置文件hdfs-site.xml内容
[root@masterNode1 ~]# vim /usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml
[root@masterNode1 ~]# #开启hadoop
[root@masterNode1 ~]# start-all.sh
This script is deprecated. Instead use start-dfs.sh and start-yarn.sh
[Fatal Error] hdfs-site.xml:54:11: The end-tag for element type "property" must end with a ">" delimiter.
23/03/24 19:00:38 PAPAI conf.Configuration: error parsing conf hdfs-site.xml
org.xml.sax.SAXParseException: systemId: file:/usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml; lineNumber: 54; columnNumber: 11; The end-tag for element type "prop
erty" must end with a ">" delimiter.
    at org.apache.xerces.parsers.DCParser.parse(Unknown Source)
    at org.apache.xerces.jaxp.DocumentBuilderImpl.parse(Unknown Source)
    at javax.xml.parsers.DocumentBuilder.parse(DocumentBuilder.java:150)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2480)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2468)
    at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:2539)
    at org.apache.hadoop.conf.Configuration.loadResources(Configuration.java:2492)
    at org.apache.hadoop.conf.Configuration.getProps(Configuration.java:2405)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1143)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1115)
    at org.apache.hadoop.conf.Configuration.setBoolean(Configuration.java:1451)
    at org.apache.hadoop.util.GenericOptionsParser.processGeneralOptions(GenericOptionsParser.java:321)
    at org.apache.hadoop.util.GenericOptionsParser.parseGeneralOptions(GenericOptionsParser.java:487)
    at org.apache.hadoop.util.GenericOptionsParser.<init>(GenericOptionsParser.java:170)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:153)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at org.apache.hadoop.hdfs.tools.GetConf.main(GetConf.java:332)
Exception in thread "main" java.lang.RuntimeException: org.xml.sax.SAXParseException; systemId: file:/usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml; lineNumber: 54; columnNumber: 11; The end-tag for element type "property" must end with a ">" delimiter.
    at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:2492)
    at org.apache.hadoop.conf.Configuration.loadResources(Configuration.java:2405)
    at org.apache.hadoop.conf.Configuration.getProps(Configuration.java:2405)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1143)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1115)
    at org.apache.hadoop.conf.Configuration.setBoolean(Configuration.java:1451)
    at org.apache.hadoop.util.GenericOptionsParser.processGeneralOptions(GenericOptionsParser.java:321)
    at org.apache.hadoop.util.GenericOptionsParser.parseGeneralOptions(GenericOptionsParser.java:487)
    at org.apache.hadoop.util.GenericOptionsParser.<init>(GenericOptionsParser.java:170)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:153)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at org.apache.hadoop.hdfs.tools.GetConf.main(GetConf.java:332)
Caused by: org.xml.sax.SAXParseException; systemId: file:/usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml; lineNumber: 54; columnNumber: 11; The end-tag for element
type "property" must end with a ">" delimiter.
    at org.apache.xerces.parsers.DCParser.parse(Unknown Source)
    at org.apache.xerces.jaxp.DocumentBuilderImpl.parse(Unknown Source)
    at javax.xml.parsers.DocumentBuilder.parse(DocumentBuilder.java:150)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2480)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2468)
    at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:2539)
    ... 12 more
Starting namenodes on [ ]
masterNode1: Warning: Permanently added 'masterNode1,100.100.104.48' (ECDSA) to the list of known hosts.
slaveNode2: Warning: Permanently added 'slaveNode2,100.100.105.20' (ECDSA) to the list of known hosts.
slaveNode1: Warning: Permanently added 'slaveNode1,100.100.106.13' (ECDSA) to the list of known hosts.
masterNode1: starting namenode, logging to /usr/hadoop-2.7.1/logs/hadoop-root-namenode-masterNode1.out
slaveNode1: starting namenode, logging to /usr/hadoop-2.7.1/logs/hadoop-root-namenode-slaveNode2.out
slaveNode1: [Fatal Error] hdfs-site.xml:54:11: The end-tag for element type "property" must end with a ">" delimiter.
masterNode1: Warning: Permanently added 'masterNode1,100.100.104.48' (ECDSA) to the list of known hosts.
slaveNode2: Warning: Permanently added 'slaveNode2,100.100.105.20' (ECDSA) to the list of known hosts.
slaveNode1: Warning: Permanently added 'slaveNode1,100.100.106.13' (ECDSA) to the list of known hosts.
masterNode1: starting datanode, logging to /usr/hadoop-2.7.1/logs/hadoop-root-datanode-masterNode1.out
slaveNode1: starting datanode, logging to /usr/hadoop-2.7.1/logs/hadoop-root-datanode-slaveNode1.out
slaveNode1: starting datanode, logging to /usr/hadoop-2.7.1/logs/hadoop-root-datanode-slaveNode2.out
masterNode1: [Fatal Error] hdfs-site.xml:54:11: The end-tag for element type "property" must end with a ">" delimiter.
[Fatal Error] hdfs-site.xml:54:11: The end-tag for element type "property" must end with a ">" delimiter.
23/03/24 19:08:38 PAPAI conf.Configuration: error parsing conf hdfs-site.xml
org.xml.sax.SAXParseException: systemId: file:/usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml; lineNumber: 54; columnNumber: 11; The end-tag for element type "prop
erty" must end with a ">" delimiter.
    at org.apache.xerces.parsers.DCParser.parse(Unknown Source)
    at org.apache.xerces.jaxp.DocumentBuilderImpl.parse(Unknown Source)
    at javax.xml.parsers.DocumentBuilder.parse(DocumentBuilder.java:150)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2480)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2468)
    at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:2539)
    at org.apache.hadoop.conf.Configuration.loadResources(Configuration.java:2492)
    at org.apache.hadoop.conf.Configuration.getProps(Configuration.java:2405)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1143)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1115)
    at org.apache.hadoop.conf.Configuration.setBoolean(Configuration.java:1451)
    at org.apache.hadoop.util.GenericOptionsParser.processGeneralOptions(GenericOptionsParser.java:321)
    at org.apache.hadoop.util.GenericOptionsParser.parseGeneralOptions(GenericOptionsParser.java:487)
    at org.apache.hadoop.util.GenericOptionsParser.<init>(GenericOptionsParser.java:170)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:153)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at org.apache.hadoop.hdfs.tools.GetConf.main(GetConf.java:332)
Exception in thread "main" java.lang.RuntimeException: org.xml.sax.SAXParseException; systemId: file:/usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml; lineNumber: 54; columnNumber: 11; The end-tag for element type "property" must end with a ">" delimiter.
    at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:2492)
    at org.apache.hadoop.conf.Configuration.loadResources(Configuration.java:2405)
    at org.apache.hadoop.conf.Configuration.getProps(Configuration.java:2405)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1143)
    at org.apache.hadoop.conf.Configuration.set(Configuration.java:1115)
    at org.apache.hadoop.conf.Configuration.setBoolean(Configuration.java:1451)
    at org.apache.hadoop.util.GenericOptionsParser.processGeneralOptions(GenericOptionsParser.java:321)
    at org.apache.hadoop.util.GenericOptionsParser.parseGeneralOptions(GenericOptionsParser.java:487)
    at org.apache.hadoop.util.GenericOptionsParser.<init>(GenericOptionsParser.java:170)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:153)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at org.apache.hadoop.hdfs.tools.GetConf.main(GetConf.java:332)
Caused by: org.xml.sax.SAXParseException; systemId: file:/usr/hadoop-2.7.1/etc/hadoop/hdfs-site.xml; lineNumber: 54; columnNumber: 11; The end-tag for element
type "property" must end with a ">" delimiter.
    at org.apache.xerces.parsers.DCParser.parse(Unknown Source)
    at org.apache.xerces.jaxp.DocumentBuilderImpl.parse(Unknown Source)
    at javax.xml.parsers.DocumentBuilder.parse(DocumentBuilder.java:150)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2480)
    at org.apache.hadoop.conf.Configuration.parse(Configuration.java:2468)
    at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:2539)
    ... 12 more
Starting yarn daemons
starting resourcemanager, logging to /usr/hadoop-2.7.1/logs/yarn-root-resourcemanager-masterNode1.out
masterNode1: Warning: Permanently added 'masterNode1,100.100.104.48' (ECDSA) to the list of known hosts.
slaveNode1: Warning: Permanently added 'slaveNode1,100.100.106.13' (ECDSA) to the list of known hosts.
slaveNode2: Warning: Permanently added 'slaveNode2,100.100.105.20' (ECDSA) to the list of known hosts.
masterNode1: starting nodemanager, logging to /usr/hadoop-2.7.1/logs/yarn-root-nodemanager-masterNode1.out
slaveNode1: starting nodemanager, logging to /usr/hadoop-2.7.1/logs/yarn-root-nodemanager-slaveNode1.out
slaveNode2: starting nodemanager, logging to /usr/hadoop-2.7.1/logs/yarn-root-nodemanager-slaveNode2.out
[root@masterNode1 ~]# #查看acl
```

## 创建 aclFile.txt 文件、设定文件权限

```
[root@masterNode1 ~]# hdfs dfs -getfacl -R /file
# file: /file
# owner: root
# group: root
user::rwx
group::r-x
other::r-x

[root@masterNode1 ~]# #创建 aclFile.txt文件
[root@masterNode1 ~]# hdfs dfs -touchz /datafile/aclFile.txt
[root@masterNode1 ~]# #设定文件权限
[root@masterNode1 ~]# #设定 -m: 修改acl
[root@masterNode1 ~]# hdfs dfs -setfacl -m user:hadoop:rw- /datafile/aclFile.txt
[root@masterNode1 ~]# hdfs dfs -getfacl /datafile/aclFile.txt
# file: /datafile/aclFile.txt
# owner: root
# group: supergroup
user::rw-
user:hadoop:rw-
group::r--
mask::rw-
other::r--

[root@masterNode1 ~]# #设定 -x: 删除指定规则
[root@masterNode1 ~]# hdfs dfs -setfacl -x user:hadoop /datafile/aclFile.txt
[root@masterNode1 ~]# hdfs dfs -getfacl /datafile/aclFile.txt
# file: /datafile/aclFile.txt
# owner: root
# group: supergroup
user::rw-
group::r--
mask::r--
other::r--

[root@masterNode1 ~]# #设定 -b: 基本的acl规则(所有者, 群组, 其他)被保留, 其他规则全部删除
[root@masterNode1 ~]# hdfs dfs -setfacl -b /datafile/aclFile.txt
[root@masterNode1 ~]# hdfs dfs -getfacl /datafile/aclFile.txt
# file: /datafile/aclFile.txt
# owner: root
# group: supergroup
user::rw-
group::r--
other::r--

[root@masterNode1 ~]# #设定 -k: 删除默认的ACL
[root@masterNode1 ~]# hdfs dfs -setfacl -k /datafile/aclFile.txt
[root@masterNode1 ~]# hdfs dfs -getfacl /datafile/aclFile.txt
# file: /datafile/aclFile.txt
# owner: root
# group: supergroup
user::rw-
group::r--
other::r--

[root@masterNode1 ~]#
```

```
Last login: Fri Jan 8 16:10:24 2021 from 192.168.1.71
[root@masterNode1 ~]# 集群服务启动完成
hdfs dfsadmin -refreshNodes
Refresh nodes successful
[root@masterNode1 ~]# hdfs dfsadmin -report -live
Configured Capacity: 71425904713728 (64.96 TB)
Present Capacity: 66428109479936 (60.42 TB)
DFS Remaining: 66428109234176 (60.42 TB)
DFS Used: 245760 (240 KB)
DFS Used%: 0.00%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0
Missing blocks (with replication factor 1): 0

-----
Live datanodes (3):

Name: 100.100.106.13:50010 (slaveNode2)
Hostname: slaveNode2
Decommission Status : Normal
Configured Capacity: 23808634216448 (21.65 TB)
DFS Used: 81920 (80 KB)
Non DFS Used: 1583460106240 (1.44 TB)
DFS Remaining: 22225174028288 (20.21 TB)
DFS Used%: 0.00%
DFS Remaining%: 93.35%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Fri Mar 24 19:27:43 CST 2023

Name: 100.100.104.48:50010 (masterNode1)
Hostname: masterNode1
Decommission Status : Normal
Configured Capacity: 23808636280832 (21.65 TB)
DFS Used: 81920 (80 KB)
Non DFS Used: 1764782809088 (1.61 TB)
DFS Remaining: 22043853389824 (20.05 TB)
DFS Used%: 0.00%
DFS Remaining%: 92.59%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Fri Mar 24 19:27:43 CST 2023

Name: 100.100.105.20:50010 (slaveNode1)
Hostname: slaveNode1
Decommission Status : Normal
Configured Capacity: 23808634216448 (21.65 TB)
DFS Used: 81920 (80 KB)
Non DFS Used: 1649552318464 (1.50 TB)
DFS Remaining: 22159081816064 (20.15 TB)
DFS Used%: 0.00%
DFS Remaining%: 93.07%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Fri Mar 24 19:27:43 CST 2023

[root@masterNode1 ~]# #在HDFS创建/quotaDir测试目录
[root@masterNode1 ~]# hdfs dfs -mkdir /quotaDir
[root@masterNode1 ~]# #为测试目录设置限制容量命令
[root@masterNode1 ~]# hdfs dfsadmin -setSpaceQuota 300000
```

```

[root@masterNode1 ~]# #在HDFS创建/quotadir测试目录
[root@masterNode1 ~]# hdfs dfs -mkdir /quotadir
[root@masterNode1 ~]# #为测试目录设置限制容量命令
[root@masterNode1 ~]# hdfs dfsadmin -setSpaceQuota 300000000 /quotadir
[root@masterNode1 ~]# #在本地创建文件并输入数据字符串如abc:
[root@masterNode1 ~]# vim /usr/test.txt
[root@masterNode1 ~]# #将test.txt文件上传至HDFS测试目录中（成功）
[root@masterNode1 ~]# hdfs dfs -put /usr/test.txt /quotadir/text.txt
[root@masterNode1 ~]# #移除测试目录的容量
[root@masterNode1 ~]# hdfs dfsadmin -setSpaceQuota 30 /quotadir
[root@masterNode1 ~]# #继续将test.txt文件上传至HDFS测试目录中（失败）
[root@masterNode1 ~]# hdfs dfs -put /usr/test.txt /quotadir/text1.txt
23/03/24 15:21:11 WARN hdfs.NFSClient: DataStreamException
org.apache.hadoop.hdfs.protocol.DSQuotaExceededException: The DiskSpace quota of /quotadir is exceeded: quota = 30 B = 30 B but diskpace consumed = 67108890 B = 64.00 MB
    at org.apache.hadoop.hdfs.server.namenode.DirectoryWithQuotaFeature.verifyStorageSpaceQuota(DirectoryWithQuotaFeature.java:211)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.verifyQuota(FSDirectory.java:866)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.updateCount(FSDirectory.java:699)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.updateCount(FSDirectory.java:658)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.addBlock(FSDirectory.java:483)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.saveAllocatedBlock(FSNamesystem.java:3574)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.storeAllocatedBlock(FSNamesystem.java:3160)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.getAdditionalBlock(FSNamesystem.java:3041)
    at org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.addBlock(NameNodeRpcServer.java:723)
    at org.apache.hadoop.hdfs.protocol.proto.ClientNameNodeProtocolServerSideTranslatorPB.addBlock(ClientNameNodeProtocolServerSideTranslatorPB.java:492)
    at org.apache.hadoop.hdfs.protocol.proto.ClientNameNodeProtocolProtos$ClientNameNodeProtocol$2.callBlockingMethod(ClientNameNodeProtocolProtos.java)
    at org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine.java:616)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:969)
    at org.apache.hadoop.ipc.Server$Handler$1.run(Server.java:2049)
    at org.apache.hadoop.ipc.Server$Handler$1.run(Server.java:2045)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1657)
    at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2043)

    at sun.reflect.NativeConstructorAccessorImpl.newInstance(Native Method)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)
    at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
    at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
    at org.apache.hadoop.ipc.RemoteException.instantiateException(RemoteException.java:106)
    at org.apache.hadoop.ipc.RemoteException.unwrapRemoteException(RemoteException.java:73)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.locateFollowingBlock(DFSOutputStream.java:1433)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.nextBlockOutputStream(DFSOutputStream.java:1226)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:449)
Caused by: org.apache.hadoop.ipc.RemoteException(org.apache.hadoop.hdfs.protocol.DSQuotaExceededException): The DiskSpace quota of /quotadir is exceeded: quota = 30 B = 30 B but diskpace consumed = 67108890 B = 64.00 MB
    at org.apache.hadoop.hdfs.server.namenode.DirectoryWithQuotaFeature.verifyStorageSpaceQuota(DirectoryWithQuotaFeature.java:211)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.verifyQuota(FSDirectory.java:866)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.updateCount(FSDirectory.java:699)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.updateCount(FSDirectory.java:658)
    at org.apache.hadoop.hdfs.server.namenode.FSDirectory.addBlock(FSDirectory.java:483)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.saveAllocatedBlock(FSNamesystem.java:3574)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.storeAllocatedBlock(FSNamesystem.java:3160)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.getAdditionalBlock(FSNamesystem.java:3041)
    at org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.addBlock(NameNodeRpcServer.java:723)
    at org.apache.hadoop.hdfs.protocol.proto.ClientNameNodeProtocolServerSideTranslatorPB.addBlock(ClientNameNodeProtocolServerSideTranslatorPB.java:492)
    at org.apache.hadoop.hdfs.protocol.proto.ClientNameNodeProtocolProtos$ClientNameNodeProtocol$2.callBlockingMethod(ClientNameNodeProtocolProtos.java)
    at org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(ProtobufRpcEngine.java:616)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:969)
    at org.apache.hadoop.ipc.Server$Handler$1.run(Server.java:2049)
    at org.apache.hadoop.ipc.Server$Handler$1.run(Server.java:2045)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1657)
    at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2043)

    at org.apache.hadoop.ipc.Client.call(Client.java:1476)
    at org.apache.hadoop.ipc.Client.call(Client.java:1407)
    at org.apache.hadoop.ipc.ProtobufRpcEngine$Invoker.invoke(ProtobufRpcEngine.java:229)
    at com.sun.proxy.$Proxy9.addBlock(Unknown Source)
    at org.apache.hadoop.hdfs.protocolPB.ClientNameNodeProtocolTranslatorPB.addBlock(ClientNameNodeProtocolTranslatorPB.java:418)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.io.retry.RetryInvocationHandler.invokeMethod(RetryInvocationHandler.java:187)
    at org.apache.hadoop.io.retry.RetryInvocationHandler.invoke(RetryInvocationHandler.java:102)
    at com.sun.proxy.$Proxy10.addBlock(Unknown Source)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.locateFollowingBlock(DFSOutputStream.java:1430)
    ... 2 more
put: The DiskSpace quota of /quotadir is exceeded: quota = 30 B = 30 B but diskpace consumed = 67108890 B = 64.00 MB
[root@masterNode1 ~]# #在HDFS中创建目录
[root@masterNode1 ~]# hdfs dfs -mkdir /QuotaTest
[root@masterNode1 ~]# #设置限制文件数量命令
[root@masterNode1 ~]# hdfs dfsadmin -setQuota 2 /QuotaTest
[root@masterNode1 ~]# #在QuotaTest目录中创建文件（第一步成功）
[root@masterNode1 ~]# hdfs dfs -touchz /QuotaTest/Test.txt
[root@masterNode1 ~]# #继续在QuotaTest目录中创建文件（第二步失败）
[root@masterNode1 ~]# hdfs dfs -touchz /QuotaTest/Test1.txt
touchz: The NameSpace quota (directories and files) of directory /QuotaTest is exceeded: quota=2 file count=3
[root@masterNode1 ~]# #查看quotadir目录的信息
[root@masterNode1 ~]# hdfs dfs -count -q /quotadir
      none      inf      30      4      1      1      13 /quotadir
[root@masterNode1 ~]# #查看QuotaTest目录的信息
[root@masterNode1 ~]# hdfs dfs -count -q /QuotaTest
      2      0      none      inf      1      1      0 /QuotaTest

```

## 在 HDFS 中创建目录

```

[root@masterNode1 ~]# #在HDFS中创建目录
[root@masterNode1 ~]# hdfs dfs -mkdir /QuotaTest
[root@masterNode1 ~]# #设置限制文件数量命令
[root@masterNode1 ~]# hdfs dfsadmin -setQuota 2 /QuotaTest
[root@masterNode1 ~]# #在QuotaTest目录中创建文件（第一步成功）
[root@masterNode1 ~]# hdfs dfs -touchz /QuotaTest/Test.txt
[root@masterNode1 ~]# #继续在QuotaTest目录中创建文件（第二步失败）
[root@masterNode1 ~]# hdfs dfs -touchz /QuotaTest/Test1.txt
touchz: The NameSpace quota (directories and files) of directory /QuotaTest is exceeded: quota=2 file count=3
[root@masterNode1 ~]# #查看quotadir目录的信息
[root@masterNode1 ~]# hdfs dfs -count -q /quotadir
      none      inf      30      4      1      1      13 /quotadir
[root@masterNode1 ~]# #查看QuotaTest目录的信息
[root@masterNode1 ~]# hdfs dfs -count -q /QuotaTest
      2      0      none      inf      1      1      0 /QuotaTest

```

## 清除空间配额、检查集群文件

```
[root@masterNode1 ~]# #清除空间配额
[root@masterNode1 ~]# hdfs dfsadmin -clrSpaceQuota /quotaDir
[root@masterNode1 ~]# #继续将test.txt文件上传至HDFS测试目录中
[root@masterNode1 ~]# hdfs dfs -put /usr/test.txt /quotaDir/text1.txt
[root@masterNode1 ~]# #清除文件数量限制配额
[root@masterNode1 ~]# hdfs dfsadmin -clrQuota /QuotaTest
[root@masterNode1 ~]# #继续在QuotaTest目录中创建文件
[root@masterNode1 ~]# hdfs dfs -touchz /QuotaTest/Test1.txt
[root@masterNode1 ~]#
[root@masterNode1 ~]# #创建/fsck目录
[root@masterNode1 ~]# hdfs dfs -mkdir /fsck
[root@masterNode1 ~]# #检查集群文件
[root@masterNode1 ~]# hdfs fsck /fsck -files -blocks -locations
Connecting to namenode via http://masterNode1:50070/fsck?ugi=root&files=1&blocks=1&locations=1&path=%2Ffsck
FSCK started by root (auth:SIMPLE) from /100.100.104.48 for path /fsck at Fri Mar 24 19:36:57 CST 2023
/fsck <dir>
Status: HEALTHY
Total size:      0 B
Total dirs:      1
Total files:      0
Total symlinks:      0
Total blocks (validated):      0
Minimally replicated blocks:    0
Over-replicated blocks:        0
Under-replicated blocks:       0
Mis-replicated blocks:         0
Default replication factor:     2
Average block replication:      0.0
Corrupt blocks:                 0
Missing replicas:               0
Number of data-nodes:          3
Number of racks:                1
FSCK ended at Fri Mar 24 19:36:57 CST 2023 in 2 milliseconds

The filesystem under path '/fsck' is HEALTHY
```

```
[root@masterNode1 ~]# #设置允许的数据利用率阈值
[root@masterNode1 ~]# start-balancer.sh -threshold 10
starting balancer, logging to /usr/hadoop-2.7.1/logs/hadoop-root-balancer-masterNode1.out
Time Stamp      Iteration#  Bytes Already Moved  Bytes Left To Move  Bytes Being Moved
[root@masterNode1 ~]# #设置带宽为1048576
[root@masterNode1 ~]# hdfs dfsadmin -setBalancerBandwidth 1048576
Balancer bandwidth is set to 1048576
[root@masterNode1 ~]# #在masterNode1上执行
[root@masterNode1 ~]# hdfs balancer -include slaveNode1,slaveNode2

23/03/24 19:42:36 INFO balancer.Balancer: namenodes = [hdfs://masterNode1:9000]
23/03/24 19:42:36 INFO balancer.Balancer: parameters = Balancer.Parameters[BalancingPolicy.Node, threshold=10.0, max idle iteration = 5, number of nodes to be excluded = 0, number of nodes to be included = 2]
Time Stamp      Iteration#  Bytes Already Moved  Bytes Left To Move  Bytes Being Moved
23/03/24 19:42:37 INFO net.NetworkTopology: Adding a new node: /default-rack/100.100.106.13:50010
23/03/24 19:42:37 INFO net.NetworkTopology: Adding a new node: /default-rack/100.100.105.20:50010
23/03/24 19:42:37 INFO balancer.Balancer: 0 over-utilized: []
23/03/24 19:42:37 INFO balancer.Balancer: 0 underutilized: []
The cluster is balanced. Exiting...
2023-3-24 19:42:37      0      0 B      0 B      -1 B
2023-3-24 19:42:37      Balancing took 804.0 milliseconds
```

## 第二次作业

进入 shell 模式

```
<configuration>
<property>
<name>hbase.rootdir</name>
<value>hdfs://masterNode1:9000/hbase</value>
<description>指定node1的hdfs协议的文件系统通信地址，这个需要和Hadoop的端口一致</description>
</property>

<property>
<name>hbase.cluster.distributed</name>
<value>true</value>
<description>开启hbase集群</description>
</property>

<property>
<name>hbase.master</name>[]
<value>masterNode1:60000</value>
<description>指定主机ID和地址</description>
</property>

<property>
<name>hbase.master.info.port</name>
<value>60010</value>
<description>配置hbaseweb访问端口</description>
</property>

<property>
<name>hbase.tmp.dir</name>
<value>/home/hbase/tmp</value>
<description>hbase临时数据的目录</description>
</property>

<property>
<name>hbase.zookeeper.quorum</name>
<value>masterNode1,slaveNode1,slaveNode2</value>
<description>zookeeper集群节点</description>
</property>

<property>
<name>hbase.zookeeper.property.clientPort</name>
<value>2181</value>
<description>zookeeper的访问端口</description>
</property>
</configuration>
~
```

```
# see http://wiki.apache.org/hadoop/PerformanceTuning
export HBASE_OPTS="-XX:+UseConcMarkSweepGC"
export JAVA_HOME=/usr/java/jdk1.8.0_162
export HBASE_MANAGES_ZK=false[]
#Configure PermSize. Only needed in JDK7. You can safely remo
```

```
localhost
masterNode1
slaveNode1
slaveNode2[]
~
~
```

```
Last login: Fri Apr 7 18:30:38 2023 from 189.80.43.11
[root@masterNode1 ~]# /usr/hbase-1.3.2/bin/hbase shell
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hbase-1.3.2/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hadoop-2.7.1/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Hbase Shell: enter "help<RETURN>" for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.3.2, r1bedb5bfb5a99067e7bc54718c3124f632b6e17, Mon Mar 19 18:47:19 UTC 2018
```



## 查看 HBase 状态

```
hbase(main):001:0> status
1 active master, 0 backup masters, 1 servers, 2 dead, 2.0000 average load
```

## 查看 HBase 版本信息

```
hbase(main):002:0> version
1.3.2, r1bedb5bfbb5a99067e7bc54718c3124f632b6e17, Mon Mar 19 18:47:19 UTC 2018
```

## 创建 user 表

```
hbase(main):003:0> create 'user','address','info' #创建user表
0 row(s) in 1.3230 seconds

=> Hbase::Table - user
```

## 描述 user 表

```
hbase(main):004:0> describe 'user' #描述user表
Table user is ENABLED
user
COLUMN FAMILIES DESCRIPTION
(NAME => 'address', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER',
, COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
(NAME => 'info', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER',
, COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
2 row(s) in 0.1030 seconds
```

## 查询表

```
hbase(main):005:0> list #查询表
TABLE
user
1 row(s) in 0.0220 seconds

=> ["user"]
```

## 插入数据

```
=> ["user"]
hbase(main):006:0> put 'user','xiaoming','address:home','GuangDong'
0 row(s) in 0.1100 seconds

hbase(main):007:0> put 'user','xiaoming','address:school','daxue'
0 row(s) in 0.0190 seconds

hbase(main):008:0> put 'user','xiaoming','info:sex','man'
0 row(s) in 0.0210 seconds

hbase(main):009:0> put 'user','xiaoming','info:age',18
0 row(s) in 0.0190 seconds

hbase(main):010:0> put 'user','lihong','address:home','GuangXi'
0 row(s) in 0.0180 seconds

hbase(main):011:0> put 'user','lihong','address:school','zhongxue'
0 row(s) in 0.0180 seconds

hbase(main):012:0> put 'user','lihong','info:sex','woman'
0 row(s) in 0.0190 seconds

hbase(main):013:0> put 'user','lihong','info:age',18
0 row(s) in 0.0190 seconds
```

## 查询数据

```
hbase(main):014:0> scan 'user'
ROW                                COLUMN+CELL
lihong                            column=address:home, timestamp=1680863222497, value=GuangXi
lihong                            column=address:school, timestamp=1680863228500, value=zhongxue
lihong                            column=info:age, timestamp=1680863240028, value=18
lihong                            column=info:sex, timestamp=1680863234067, value=woman
xiaoming                          column=address:home, timestamp=1680863196970, value=GuangDong
xiaoming                          column=address:school, timestamp=1680863203808, value=daxue
xiaoming                          column=info:age, timestamp=1680863216767, value=18
xiaoming                          column=info:sex, timestamp=1680863210827, value=man
2 row(s) in 0.0540 seconds
```

```
hbase(main):015:0> scan 'user', FILTER=>"RowFilter(=,'binary:lihong')"
```

ROW	COLUMN+CELL
lihong	column=address:home, timestamp=1680863222497, value=GuangXi
lihong	column=address:school, timestamp=1680863228500, value=zhongxue
lihong	column=info:age, timestamp=1680863240028, value=18
lihong	column=info:sex, timestamp=1680863234067, value=woman

```
1 row(s) in 0.0730 seconds
```

```
hbase(main):016:0> get 'user','lihong'
```

COLUMN	CELL
address:home	timestamp=1680863222497, value=GuangXi
address:school	timestamp=1680863228500, value=zhongxue
info:age	timestamp=1680863240028, value=18
info:sex	timestamp=1680863234067, value=woman

```
1 row(s) in 0.0300 seconds
```

```
hbase(main):017:0> get 'user','xiaoming','info:age'
```

COLUMN	CELL
info:age	timestamp=1680863216767, value=18

```
1 row(s) in 0.0270 seconds
```

## 查询任何一列中包含 “ x ” 的数据

```
hbase(main):018:0> scan 'user', FILTER=>"RowFilter(=,'substring:i')"
```

ROW	COLUMN+CELL
lihong	column=address:home, timestamp=1680863222497, value=GuangXi
lihong	column=address:school, timestamp=1680863228500, value=zhongxue
lihong	column=info:age, timestamp=1680863240028, value=18
lihong	column=info:sex, timestamp=1680863234067, value=woman
xiaoming	column=address:home, timestamp=1680863196970, value=GuangDong
xiaoming	column=address:school, timestamp=1680863203808, value=daxue
xiaoming	column=info:age, timestamp=1680863216767, value=18
xiaoming	column=info:sex, timestamp=1680863210827, value=man

```
2 row(s) in 0.0440 seconds
```

```
hbase(main):019:0> scan 'user', FILTER=>"ValueFilter(=,'binary:GuangDong')"
```

ROW	COLUMN+CELL
xiaoming	column=address:home, timestamp=1680863196970, value=GuangDong

```
1 row(s) in 0.0450 seconds
```

```
hbase(main):020:0> scan 'user', FILTER=>"ValueFilter(=,'substring:x')"
```

ROW	COLUMN+CELL
lihong	column=address:home, timestamp=1680863222497, value=GuangXi
lihong	column=address:school, timestamp=1680863228500, value=zhongxue
xiaoming	column=address:school, timestamp=1680863203808, value=daxue

```
2 row(s) in 0.0260 seconds
```

```
hbase(main):021:0> scan 'user', FILTER=>"ColumnPrefixFilter('age') AND ValueFilter(=,'binary:18')"
```

ROW	COLUMN+CELL
lihong	column=info:age, timestamp=1680863240028, value=18
xiaoming	column=info:age, timestamp=1680863216767, value=18

```
2 row(s) in 0.0510 seconds
```

```
hbase(main):022:0> count 'user'
```

```
2 row(s) in 0.0320 seconds
```

```
=> 2
```

```
hbase(main):030:0> put 'user','lihong','message:father','hello,my son'
```

```
0 row(s) in 0.0240 seconds
```

```
hbase(main):031:0> get 'user','lihong'
```

COLUMN	CELL
address:home	timestamp=1680863222497, value=GuangXi
address:school	timestamp=1680863228500, value=zhongxue
info:age	timestamp=1680863367953, value=20
info:sex	timestamp=1680863234067, value=woman
message:father	timestamp=1680863494667, value=hello,my son

```
1 row(s) in 0.0290 seconds
```

## 删除列族

```
hbase(main):032:0> #删除列族
hbase(main):033:0> disable 'user'
0 row(s) in 2.2850 seconds

hbase(main):034:0> alter 'user','delete'=>'message'
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.7180 seconds

hbase(main):035:0> enable 'user'
0 row(s) in 1.2900 seconds

hbase(main):036:0> describe 'user'
Table user is ENABLED
user
COLUMN FAMILIES DESCRIPTION
(NAME => 'address', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
(NAME => 'info', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
2 row(s) in 0.0160 seconds
```

## 数据版本控制

```
hbase(main):043:0> #数据版本控制
hbase(main):044:0> #修改版本保存数
hbase(main):045:0> describe 'user'
Table user is ENABLED
user
COLUMN FAMILIES DESCRIPTION
(NAME => 'address', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
(NAME => 'info', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
2 row(s) in 0.0370 seconds

hbase(main):046:0> alter 'user',NAME=>'address',VERSIONS=>3
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.7000 seconds

hbase(main):047:0> alter 'user',NAME=>'info',VERSIONS=>3
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.4370 seconds

hbase(main):048:0> describe 'user'
Table user is ENABLED
user
COLUMN FAMILIES DESCRIPTION
(NAME => 'address', BLOOMFILTER => 'ROW', VERSIONS => '3', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
(NAME => 'info', BLOOMFILTER => 'ROW', VERSIONS => '3', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0')
2 row(s) in 0.0280 seconds
```

## 使用版本号

```
hbase(main):049:0> #使用版本号
hbase(main):050:0> put 'user','xiaoming','address:home','beijing'
0 row(s) in 0.1530 seconds

hbase(main):051:0> put 'user','xiaoming','address:home','shanghai'
0 row(s) in 0.0210 seconds

hbase(main):052:0> put 'user','xiaoming','address:home','wuhan'
0 row(s) in 0.0240 seconds

hbase(main):053:0> get 'user','xiaoming'
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
1 row(s) in 0.0120 seconds

hbase(main):054:0> get 'user','xiaoming',{COLUMN=>'address',VERSIONS=>3}
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
address:home                          timestamp=1680864038277, value=shanghai
address:home                          timestamp=1680864030405, value=beijing
1 row(s) in 0.0180 seconds
```

## 更新数据

```
hbase(main):055:0> put 'user','xiaoming','address:home','nanjing',1558608881450
0 row(s) in 0.0240 seconds

hbase(main):056:0> get 'user','xiaoming',{COLUMN=>'address',VERSIONS=>3}
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
address:home                          timestamp=1680864038277, value=shanghai
address:home                          timestamp=1680864030405, value=beijing
1 row(s) in 0.0140 seconds

hbase(main):057:0> put 'user','xiaoming','address:home','nanjing',1558608881450
0 row(s) in 0.0280 seconds

hbase(main):058:0> get 'user','xiaoming'
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
1 row(s) in 0.0160 seconds

hbase(main):059:0> get 'user','xiaoming',{COLUMN=>'address',VERSIONS=>3}
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
address:home                          timestamp=1680864038277, value=shanghai
address:home                          timestamp=1680864030405, value=beijing
1 row(s) in 0.0150 seconds
```

```
hbase(main):057:0> put 'user','xiaoming','address:home','nanjing',1558608881450
0 row(s) in 0.0280 seconds

hbase(main):058:0> get 'user','xiaoming'
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
1 row(s) in 0.0160 seconds

hbase(main):059:0> get 'user','xiaoming',{COLUMN=>'address',VERSIONS=>3}
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
address:home                          timestamp=1680864038277, value=shanghai
address:home                          timestamp=1680864030405, value=beijing
1 row(s) in 0.0150 seconds

hbase(main):060:0> put 'user','xiaoming','address:home','nanjing',1680864038277
0 row(s) in 0.0180 seconds

hbase(main):061:0> get 'user','xiaoming',{COLUMN=>'address',VERSIONS=>3}
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
address:home                          timestamp=1680864038277, value=nanjing
address:home                          timestamp=1680864030405, value=beijing
1 row(s) in 0.0150 seconds
```

## 读取最大 VERSIONS 个记录

```
hbase(main):001:0> delete 'user','xiaoming','address:home', 1680864038277
0 row(s) in 0.2490 seconds

hbase(main):002:0> get 'user','xiaoming',{COLUMN=>'address',VERSIONS=>3}
COLUMN                                CELL
address:home                          timestamp=1680864044828, value=wuhan
address:home                          timestamp=1680864030405, value=beijing
address:home                          timestamp=1558608881450, value=nanjing
1 row(s) in 0.0360 seconds
```

## 布隆过滤器

```
hbase(main):012:0> #布隆过滤器
hbase(main):013:0* alter 'user',NAME=>'xiaoming',BLOOMFILTER=>'ROW'
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.6190 seconds

hbase(main):014:0> alter 'user',NAME=>'xiaoming',BLOOMFILTER=>'ROWCOL'
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.4660 seconds

hbase(main):015:0> []
```

```
<configuration>
<property>
<name>hbase.rootdir</name>
<value>hdfs://masterNode1:9000/hbase</value>
<description>指定node1的hdfs协议的文件系统通信地址，这个需要和Hadoop的端口一致</description>
</property>

<property>
<name>hbase.cluster.distributed</name>
<value>true</value>
<description>开启hbase集群</description>
</property>

<property>
<name>hbase.master</name>
<value>masterNode1:60000</value>
<description>指定主机id和地址</description>
</property>

<property>
<name>hbase.master.info.port</name>
<value>60010</value>
<description>配置hbaseweb访问端口</description>
</property>

<property>
<name>hbase.tmp.dir</name>
<value>/home/hbase/tmp</value>
<description>hbase临时数据的目录</description>
</property>

<property>
<name>hbase.zookeeper.quorum</name>
<value>masterNode1,slaveNode1,slaveNode2</value>
<description>zookeeper集群节点</description>
</property>

<property>
<name>hbase.zookeeper.property.clientPort</name>
<value>2181</value>
<description>zookeeper的访问端口</description>
</property>
</configuration>
```

## 修改 JAVA\_HOME 路径

```
[root@slaveNode2 ~]# #修改JAVA_HOME 路径
[root@slaveNode2 ~]# cat /etc/profile
# /etc/profile

# System wide environment and startup programs, for login setup
# Functions and aliases go in /etc/bashrc

# It's NOT a good idea to change this file unless you know what you
# are doing. It's much better to create a custom.sh shell script in
# /etc/profile.d/ to make custom changes to your environment, as this
# will prevent the need for merging in future updates.

pathmunge () {
    case "${PATH}:" in
        *:"$1":*)
            ;;
        *)
            if [ "$2" = "after" ] ; then
                PATH=$PATH:$1
            else
                PATH=$1:$PATH
            fi
        esac
    }

if [ -x /usr/bin/id ]; then
    if [ -z "$EUID" ]; then
        # ksh workaround
        EUID=`/usr/bin/id -u`
        UID=`/usr/bin/id -ru`
    fi
    USER=`/usr/bin/id -un`
    LOGNAME=$USER
    MAIL="/var/spool/mail/$USER"
fi

# Path manipulation
if [ "$EUID" = "0" ]; then
    pathmunge /usr/sbin
    pathmunge /usr/local/sbin
else
    pathmunge /usr/local/sbin after
    pathmunge /usr/sbin after
fi

HOSTNAME=`/usr/bin/hostname 2>/dev/null`
HISTSIZE=1000
if [ "$HISTCONTROL" = "ignorespace" ] ; then
    export HISTCONTROL=ignoreboth
else
    export HISTCONTROL=ignoredups
fi

export PATH USER LOGNAME MAIL HOSTNAME HISTSIZE HISTCONTROL

# By default, we want umask to get set. This sets it for login shell
# Current threshold for system reserved uid/gids is 200
# You could check uidgid reservation validity in
# /usr/share/doc/setup-*/uidgid file
if [ $UID -gt 199 ] && [ "`/usr/bin/id -gn`" = "`/usr/bin/id -un`" ]; then
    umask 002
else
    umask 022
fi

for i in /etc/profile.d/*.sh /etc/profile.d/sh.local ; do
    if [ -r "$i" ]; then
        if [ "${#i}" != "$-" ]; then
            . "$i"
        else
            . "$i" >/dev/null
        fi
    fi
done

unset i
unset -f pathmunge
export LC_ALL=zh_CN.utf8
export JAVA_HOME=/usr/java/jdk1.8.0_162
export CLASSPATH=$JAVA_HOME/lib/dt.jar:$JAVA_HOME/lib/tools.jar:$JRE_HOME
export PATH=$PATH:$JAVA_HOME/bin:$JRE_HOME/bin:$CLASSPATH
export PATH=$PATH:$JAVA_HOME/bin
export HADOOP_HOME=/usr/hadoop-2.7.1
export PATH=$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
export SCALA_HOME=/usr/scala
export PATH=${PATH}:${SCALA_HOME}/bin
[root@slaveNode2 ~]#
```

## 修改 JKS 路径

```
# see http://wiki.apache.org/hadoop/PerformanceTuning
export HBASE_OPTS="-XX:+UseConcMarkSweepGC"
export JAVA_HOME=/usr/java/jdk1.8.0_162 #修改jdk路径
export HBASE_MANAGES_ZK=false #修改Hbase使用定义的Zookeeper，而不用内置的Zookeeper
# Configure PermSize. Only needed in JDK7. You can safely remove it for JDK8+
```

## 启动 HBase 标准方式启动

```
[root@masterNode1 ~]# #配置Hbase集群regieserver信息
[root@masterNode1 ~]# vim /usr/hbase-1.3.2/conf/regionservers #编辑regieserver文件
[root@masterNode1 ~]# #启动Hbase
[root@masterNode1 ~]# /usr/hbase-1.3.2/bin/hbase-daemon.sh start master #启动master
starting master, logging to /usr/hbase-1.3.2/bin/../logs/hbase-root-master-masterNode1.out
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option PermSize=128m; support was removed in 8.0
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=128m; support was removed in 8.0
[root@masterNode1 ~]# /usr/hbase-1.3.2/bin/hbase-daemon.sh start regionserver #启动HRegionServer
starting regionserver, logging to /usr/hbase-1.3.2/bin/../logs/hbase-root-regionserver-masterNode1.out
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option PermSize=128m; support was removed in 8.0
Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=128m; support was removed in 8.0
[root@masterNode1 ~]# /usr/java/jdk1.8.0_162/bin/jps
657 DataNode
1394 HMaster
1061 NodeManager
2037 Jps
391 NameNode
120 QuorumPeerMain
824 ResourceManager
1721 HRegionServer
[root@masterNode1 ~]# /usr/java/jdk1.8.0_162/bin/jps
657 DataNode
1394 HMaster
1061 NodeManager
391 NameNode
120 QuorumPeerMain
824 ResourceManager
1721 HRegionServer
2093 Jps
```

## 配置 HBase 集群 regieserver 信息

```
<property>

<name>hbase.cluster.distributed</name>

<value>true</value>

<description>开启hbase集群</description>

</property>

<property>

<name>hbase.master</name>

<value>60000</value>

<description>多主指定端口即可</description>

</property>
```

Backup Master: masterN...

masternode1.60010/master-status

APACHE HBASEHomeTable DetailsProceduresLocal LogsLog LevelDebug DumpMetrics DumpHBase Configuration

Backup Master

masterNode1

Current Active Master: slaveNode1

Tasks

Show All Monitored TasksShow non-RPC TasksShow All RPC Handler TasksShow Active RPC CallsShow Client OperationsView as JSON

Start Time	Description	State	Status
Fri Apr 07 19:43:39 CST 2023	Master startup	RUNNING (since 2mins, 4sec ago)	Another master is the active master, slavenode1.16000.1680868839319; waiting to become the next active master (since 2mins, 4sec ago)

Software Attributes

Attribute Name	Value	Description
HBase Version	1.3.2, revision=1bedb5bfb5a99067e7bc54718c3124f632b6e17	HBase version and revision
HBase Compiled	Mon Mar 19 18:47:19 UTC 2018, root	When HBase version was compiled and by whom
HBase Source Checksum	e7f36f244609783299ccbc41a8ee0dd4	HBase source MD5 checksum
Hadoop Version	2.5.1, revision=2e18d179e4a8065b6a9f29cf2de9451891265cce	Hadoop version and revision

It looks like you haven't started Firefox in a while. Do you want to clean it up for a fresh, like-new experience? And by the way, welcome back!

Refresh Firefox...

Backup Master: masterN...Master: slaveNode1Backup Master: slaveNo...

slavenode1.60010/master-status

APACHE HBASEHomeTable DetailsProceduresLocal LogsLog LevelDebug DumpMetrics DumpHBase Configuration

## Master slaveNode1

### Region Servers

Base StatsMemoryRequestsStorefilesCompactions				
ServerName	Start time	Version	Requests Per Second	Num. Regions
masternode1.16020.1680867474398	Fri Apr 07 19:37:54 CST 2023	1.3.2	0	2
slavenode1.16020.1680868920956	Fri Apr 07 20:02:00 CST 2023	1.3.2	0	0
slavenode2.16020.1680867486861	Fri Apr 07 19:38:06 CST 2023	1.3.2	0	0
Total:3			0	2

APACHE HBASEHomeTable DetailsProceduresLocal LogsLog LevelDebug DumpMetrics DumpHBase Configuration

## Backup Master slaveNode2

Current Active Master: slaveNode1

### Tasks

Show All Monitored TasksShow non-RPC TasksShow All RPC Handler TasksShow Active RPC CallsShow Client OperationsView as JSON			
Start Time	Description	State	Status
Fri Apr 07 19:42:18 CST 2023	Master startup	RUNNING (since 9mins, 20sec ago)	Another master is the active master, slavenode1.16000.1680868839319; waiting to become the next active master (since 8mins, 12sec ago)



```
op/mapreduce/lib/asm-5.2.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/lib/jackson-mapper-asl-1.9.13.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/lib/protobuf-java-2.5.0.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/lib/jackson-mapper-asl-1.9.13.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-common-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-jobclient-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-shuffle-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-common-2.7.1-tests.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-hs-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-core-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-app-2.7.1.jar:/usr/hadoop-2.7.1/share/hadoop/mapreduce/hadoop-mapreduce-client-hs-plugins-2.7.1.jar:/usr/hadoop-2.7.1/contrib/capacity-scheduler/*.jar
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:java.library.path=/usr/hadoop-2.7.1/lib/native
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:java.io.tmpdir=/tmp
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:java.compiler=<NA>
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:os.name=linux
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:os.arch=amd64
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:os.version=3.10.0-957.el7.x86_64
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:user.name=root
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:user.home=/root
2023-04-07 19:52:43,840 INFO [main] zookeeper.ZooKeeper: Client environment:user.dir=/root
2023-04-07 19:52:43,841 INFO [main] zookeeper.ZooKeeper: Initiating client connection, connectString=masterNode1:2181,slaveNode1:2181,slaveNode2:2181,timeout=90000,watches=org.apache.hadoop.hbase.zookeeper.PendingWatcher@13dbed9e
2023-04-07 19:52:43,862 INFO [main-SendThread(slaveNode1:2181)] zookeeper.ClientCnxn: Opening socket connection to server slaveNode1/100.100.105.3:2181. Will not attempt to authenticate using SASL (unknown error)
2023-04-07 19:52:43,872 INFO [main-SendThread(slaveNode1:2181)] zookeeper.ClientCnxn: Socket connection established to slaveNode1/100.100.105.3:2181, initiating session
2023-04-07 19:52:43,881 INFO [main-SendThread(slaveNode1:2181)] zookeeper.ClientCnxn: Session establishment complete on server slaveNode1/100.100.105.3:2181, sessionid = 0x2875b8bd10c0007, negotiated timeout = 40000
Valid region move targets:
masterNode1.16020.1680867474398
slaveNode2.16020.1680867486861
2023-04-07 19:52:44 Unloaded slaveNode1 region(s)
2023-04-07 19:52:44 Stopping regionserver on slaveNode1
slaveNode1: Warning: Permanently added 'slaveNode1,100.100.105.3' (ECDSA) to the list of known hosts.
slaveNode1: stopping regionserver.....
2023-04-07 19:52:52 Restoring balancer state to true
[root@masterNode1 ~]#
```

```
[root@slaveNode1 ~]# /usr/hbase-1.3.2/bin/hbase shell
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hbase-1.3.2/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hadoop-2.7.1/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type 'exit<RETURN>' to leave the HBase Shell
Version 1.3.2, r1bdeb5bfb5a99067e7bc54718c3124f632b6e17, Mon Mar 19 18:47:19 UTC 2018

hbase(main):001:0> balance_switch true
true
0 row(s) in 0.1920 seconds

hbase(main):002:0>
```

## 第三次作业

集群环境配置

查看 ip 地址命令

查看本地主机名

```
Last login: Fri Apr 21 17:43:13 2023 from 189.80.43.11
[root@masterNode1 ~]# #在集群各节点（masterNode1、slaveNode1、slaveNode2）执行以下命令
[root@masterNode1 ~]# ip a #查看ip地址的命令
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
5261: eth0@if5262: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:64:64:68:21 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 100.100.104.33/16 brd 100.100.255.255 scope global eth0
        valid_lft forever preferred_lft forever
[root@masterNode1 ~]# hostname #查看本地的主机名
masterNode1
[root@masterNode1 ~]#
```

slaveNode1

```
Last login: Fri Apr 21 18:01:58 2023 from 189.80.43.11
[root@slaveNode1 ~]# ip a #查看ip地址的命令
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
1985: eth0@if1986: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:64:64:69:05 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 100.100.105.5/16 brd 100.100.255.255 scope global eth0
        valid_lft forever preferred_lft forever
[root@slaveNode1 ~]# hostname #查看本地的主机名
slaveNode1
[root@slaveNode1 ~]#
```

slaveNode2

```
Last login: Tue Dec 1 19:55:21 2020 from 192.168.1.41
[root@slaveNode2 ~]# ip a #查看ip地址的命令
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
2161: eth0@if2162: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:64:64:6b:05 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 100.100.107.5/16 brd 100.100.255.255 scope global eth0
        valid_lft forever preferred_lft forever
[root@slaveNode2 ~]# hostname #查看本地的主机名
slaveNode2
[root@slaveNode2 ~]#
```

向/etc/hosts 文件中添加 IP 和 hostname，格式如下图所示，其中间使用制表符隔开。

```
127.0.0.1    localhost
::1         localhost ip6-localhost ip6-loopback
fe00::0     ip6-localnet
ff00::0     ip6-mcastprefix
ff02::1     ip6-allnodes
ff02::2     ip6-allrouters
100.100.104.33  masterNode1
100.100.105.5  slaveNode1
100.100.107.5  slaveNode2
```

分发 masterNode1 节点的 `/etc/hosts` 文件到其余节点，此时我们可以使用其余节点的 hostname 进行分发。

```
[root@masterNode1 ~]# #在masterNode1执行以下命令
[root@masterNode1 ~]#
[root@masterNode1 ~]# scp /etc/hosts slaveNode1:/etc/
Warning: Permanently added 'slavenode1,100.100.105.5' (ECDSA) to the list of known hosts.
root@slavenode1's password:
hosts
[root@masterNode1 ~]# scp /etc/hosts slaveNode2:/etc/
Warning: Permanently added 'slavenode2,100.100.107.5' (ECDSA) to the list of known hosts.
root@slavenode2's password:
hosts
[root@masterNode1 ~]#
```

让 masterNode1 节点作为 NTP 服务器的服务端。注释掉原文件所有以 ‘server’ 开头的配置项

```
# Use public servers from the pool.ntp.org project.
# Please consider joining the pool (http://www.pool.ntp.org/join.html).
#server 0.centos.pool.ntp.org iburst
#server 1.centos.pool.ntp.org iburst
#server 2.centos.pool.ntp.org iburst
#server 3.centos.pool.ntp.org iburst

#local clock, 可提供本地的服务
server ntp.sjtu.edu.cn iburst

server 127.127.1.0

fudge 127.127.1.0 stratum 8
```

```
# Use public servers from the pool.ntp.org project.
# Please consider joining the pool (http://www.pool.ntp.org/join.html).
#server 0.centos.pool.ntp.org iburst
#server 1.centos.pool.ntp.org iburst
#server 2.centos.pool.ntp.org iburst
#server 3.centos.pool.ntp.org iburst

# 指向主服务器的主机名，前提是得做主机映射才能通过主机名解析到ip
server masterNode1 iburst
```

从 yum 源安装 NTP 服务

编辑 NTP 配置文件

查看时间同步是否成功：

```
[root@masterNode1 ~]# #在各节点（masterNode1, slaveNode1, slaveNode2）执行
[root@masterNode1 ~]# /usr/sbin/ntpd
[root@masterNode1 ~]# ntpq -p
```

remote	refid	st	t	when	poll	reach	delay	offset	jitter
058176194096.ct	.INIT.	16	u	-	1024	0	0.000	0.000	0.000
ntp.telecomarme	.INIT.	16	u	-	1024	0	0.000	0.000	0.000
ntp.nic.kz	.INIT.	16	u	-	1024	0	0.000	0.000	0.000
203.113.174.44	.INIT.	16	u	-	1024	0	0.000	0.000	0.000

```
[root@slaveNode1 ~]# #在各节点（masterNode1, slaveNode1, slaveNode2）执行
[root@slaveNode1 ~]# /usr/sbin/ntpd
[root@slaveNode1 ~]# ntpq -p
```

remote	refid	st	t	when	poll	reach	delay	offset	jitter
masterNode1	.INIT.	16	u	125	128	0	0.000	0.000	0.000

```
[root@slaveNode1 ~]#
```

```
[root@slaveNode2 ~]# #在各节点（masterNode1, slaveNode1, slaveNode2）执行
[root@slaveNode2 ~]# /usr/sbin/ntpd
[root@slaveNode2 ~]# ntpq -p
```

remote	refid	st	t	when	poll	reach	delay	offset	jitter
masterNode1	.INIT.	16	u	128	128	0	0.000	0.000	0.000

## SSH 免密登录

```
Dependency Installed:
atk.x86_64 0:2.28-1.el7
cupslibs.x86_64 1:1.6.3-35.el7
fontconfig.x86_64 0:12.1.0-4.3.el7
fribidi.x86_64 0:1.0.2-1.el7
gtk-update-icon-cache.x86_64 0:3.22.30-3.el7
hicolor-icon-theme.noarch 0:0.12-7.el7
jbigkit-libs.x86_64 0:2.0-11.el7
libXcomposite.x86_64 0:0.4.4-4.1.el7
libXext.x86_64 0:1.3.3-3.el7
libXi.x86_64 0:1.7.9-1.el7
libXrender.x86_64 0:0.9.10-1.el7
libglvnd.x86_64 1:1.0.1-0.8.git5baa1e5.el7
libjpeg-turbo.x86_64 0:1.2.90-6.el7
libshai.x86_64 0:0.1.14-9.el7
libwayland-server.x86_64 0:1.15.0-1.el7
mesa-libGL.x86_64 0:18.0.5-4.el7_6
pango.x86_64 0:1.42.4-2.el7_6
avahi-libs.x86_64 0:0.6.31-19.el7
dejavu-fonts-common.noarch 0:2.33-6.el7
fontpackages-filesystem.noarch 0:1.44-8.el7
gdk-pixbuf2.x86_64 0:2.36.12-3.el7
gtk2.x86_64 0:2.24.31-1.el7
hwdm.x86_64 0:0.252-9.1.el7
libX11.x86_64 0:1.6.5-2.el7
libXcursor.x86_64 0:1.1.15-1.el7
libXfixes.x86_64 0:5.0.3-1.el7
libXinerama.x86_64 0:1.1.3-2.1.el7
libXtst.x86_64 0:1.1.4-1.el7
libXvnd-egl.x86_64 1:1.0.1-0.8.git5baa1e5.el7
libpciaccess.x86_64 0:0.14-1.el7
libtiff.x86_64 0:4.0.3-27.el7_3
libxshmfence.x86_64 0:1.2-1.el7
mesa-libgbm.x86_64 0:18.0.5-4.el7_6
pixman.x86_64 0:0.34.0-1.el7
cairo.x86_64 0:1.15.12-3.el7
dejavu-sans-fonts.noarch 0:2.33-6.el7
freetype.x86_64 0:2.8-12.el7_6
graphite2.x86_64 0:1.3.10-1.el7_3
harfbuzz.x86_64 0:1.7.5-2.el7
jasper-libs.x86_64 0:1.900.1-33.el7
libX11-common.noarch 0:1.6.5-2.el7
libXdamage.x86_64 0:1.1.4-4.1.el7
libXft.x86_64 0:2.3.2-2.el7
libXrandr.x86_64 0:1.5.1-2.el7
libdrm.x86_64 0:2.4.91-3.el7
libglvnd-glx.x86_64 1:1.0.1-0.8.git5baa1e5.el7
libpng.x86_64 2:1.5.13-7.el7_2
libwayland-client.x86_64 0:1.15.0-1.el7
mesa-libGL.x86_64 0:18.0.5-4.el7_6
mesa-libglapi.x86_64 0:18.0.5-4.el7_6

Complete!
[root@masterNode1 ~]#
```

```
StrictHostKeyChecking no
UserKnownHostsFile /dev/null

StrictHostKeyChecking=no
UserKnownHostsFile=/dev/null
```

#在 masterNode1, slaveNode1, slaveNode2 执行

`/usr/sbin/sshd`

#生成 ecdsa 密钥

`ssh-keygen -t ecdsa -f /etc/ssh/ssh_host_ecdsa_key -N ""`

#生成 ed25519 密钥

`ssh-keygen -t ed25519 -f /etc/ssh/ssh_host_ed25519_key -N ""`

#生成 rsa 密钥

`ssh-keygen -t rsa -b 2048 -f /etc/ssh/ssh_host_rsa_key`

#生成密码后，再启动一次，看是否成功，在 ps -ef 中查询一下 ssh 服务，查询 SSH 服务线程：

```
ps -ef|grep sshd
```

```
[root@masterNode1 ~]# /usr/sbin/sshd
[root@masterNode1 ~]# #生成ecdsa密钥
[root@masterNode1 ~]# ssh-keygen -t ecdsa -f /etc/ssh/ssh_host_ecdsa_key -N ""
Generating public/private ecdsa key pair.
/etc/ssh/ssh_host_ecdsa_key already exists.
Overwrite (y/n)? #生成ed25519密钥
[root@masterNode1 ~]# ssh-keygen -t ed25519 -f /etc/ssh/ssh_host_ed25519_key -N ""
Generating public/private ed25519 key pair.
/etc/ssh/ssh_host_ed25519_key already exists.
Overwrite (y/n)? y
Your identification has been saved in /etc/ssh/ssh_host_ed25519_key.
Your public key has been saved in /etc/ssh/ssh_host_ed25519_key.pub.
The key fingerprint is:
SHA256:1HW2n+0mHY5mjPGEMF80JRBc7XhJwUG4ybd0hNMaIB0 root@masterNode1
The key's randomart image is:
+---[ED25519 256]---+
|
| o+=BO=..|
| +oO= ..|
| o ++Boo=|
| . B B=o.|
| S X =o.|
| . B..+|
| o o|
|
+----[SHA256]-----+
[root@masterNode1 ~]# #生成rsa密钥
[root@masterNode1 ~]# ssh-keygen -t rsa -b 2048 -f /etc/ssh/ssh_host_rsa_key
Generating public/private rsa key pair.
/etc/ssh/ssh_host_rsa_key already exists.
Overwrite (y/n)? y
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /etc/ssh/ssh_host_rsa_key.
Your public key has been saved in /etc/ssh/ssh_host_rsa_key.pub.
The key fingerprint is:
SHA256:KCacB8PeY+tmJLvTi/8u8h07KraXKx9dCQjBe1001KTU root@masterNode1
The key's randomart image is:
+---[RSA 2048]-----+
| .. .. o.oE. |
| .=. . o.o |
| .o=o . |
| oo=o . |
| ...* + . S |
| ...* o |
| ..o ooo |
| ...oO+ o |
| o++=O= |
+----[SHA256]-----+
[root@masterNode1 ~]# #生成密码后，再启动一次，看是否成功，在 ps -ef中查询一下 ssh服务，查询ssh服务线程：
[root@masterNode1 ~]# ps -ef|grep sshd
root      8      1  0 17:43 ?        00:00:00 /sbin/sshd
root      57      8  0 17:43 ?        00:00:00 sshd: root
root     101      8  0 17:44 ?        00:00:00 sshd: root
root     103      8  0 17:44 ?        00:00:00 sshd: root@pts/1
root     235    105  0 18:31 pts/1    00:00:00 grep --color=auto sshd
[root@masterNode1 ~]#
```

```
[root@slaveNode2 ~]# ssh-keygen -t rsa -f /root/.ssh/id_rsa
Generating public/private rsa key pair.
/root/.ssh/id_rsa already exists.
Overwrite (y/n)? y
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /root/.ssh/id_rsa.
Your public key has been saved in /root/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:iYv5BVqyZmXvFTOT3qVkl4hEDcoURqoVWcV6u+9D4OE root@slaveNode2
The key's randomart image is:
+---[RSA 2048]-----+
| .++o |
| .* ..o |
| o o.. . |
| o .o+. . |
| . . *+S+* o . |
| X =Eo.X + |
| B . +o+ + . |
| o . o... . |
| . .oo. |
+----[SHA256]-----+
```

将证书发送到本机及其他机器上，需要输入登录密码 `datastudio`

#在 masterNode1, slaveNode1, slaveNode2 执行

```
ssh-copy-id masterNode1
```

```
ssh-copy-id slaveNode1
```

```
ssh-copy-id slaveNode2
```

测试验证，使用在各节点中使用 SSH 命令分别进行登录测试，查看能否无密码登入，均可登录即设置成功，可使用 `exit` 命令退出。

```
[root@masterNode1 ~]# ssh-copy-id masterNode1
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/root/.ssh/id_rsa.pub"
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to install the new keys
Warning: Permanently added 'masterNode1,100.100.104.33' (ECDSA) to the list of known hosts.
root@masterNode1's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'masterNode1'"
and check to make sure that only the key(s) you wanted were added.
```

```
[root@masterNode1 ~]# ssh masterNode1
Warning: Permanently added 'masterNode1,100.100.104.33' (ECDSA) to the list of known hosts.
Last failed login: Fri Apr 21 18:37:53 CST 2023 from 100.100.105.5 on ssh:notty
There were 3 failed login attempts since the last successful login.
Last login: Fri Apr 21 17:44:07 2023 from 189.80.43.11
[root@masterNode1 ~]# exit
logout
Connection to masterNode1 closed.
```

```
[root@slaveNode1 ~]# ssh slaveNode1
Warning: Permanently added 'slaveNode1,100.100.105.5' (ECDSA) to the list of known hosts.
Last failed login: Fri Apr 21 18:15:25 CST 2023 from 100.100.104.33 on ssh:notty
There was 1 failed login attempt since the last successful login.
Last login: Fri Apr 21 18:07:34 2023 from 189.80.43.11
[root@slaveNode1 ~]# exit
logout
Connection to slaveNode1 closed.
```

```
[root@slaveNode2 ~]# ssh slaveNode2
Warning: Permanently added 'slaveNode2,100.100.107.5' (ECDSA) to the list of known hosts.
Last login: Fri Apr 21 17:43:13 2023 from 189.80.43.11
[root@slaveNode2 ~]# exit
logout
Connection to slaveNode2 closed.
```

## 安装 JDK

org-netbeans-lib-profiler-ui.xml	100%	429	1.0MB/s	00:00
org-netbeans-lib-profiler-common.xml	100%	437	1.2MB/s	00:00
org-netbeans-lib-profiler.xml	100%	423	790.4KB/s	00:00
org-netbeans-modules-profiler-api.xml	100%	439	1.5MB/s	00:00
org-netbeans-modules-profiler-attach.xml	100%	445	1.6MB/s	00:00
jcconsole.jar	100%	398KB	98.5MB/s	00:00
libjawn.so	100%	7103	16.7MB/s	00:00
libjli.so	100%	102KB	67.9MB/s	00:00
dt.jar	100%	159KB	80.7MB/s	00:00
tools.jar	100%	17MB	111.4MB/s	00:00
packager.jar	100%	4646	11.9MB/s	00:00
ea-jdi.jar	100%	2359KB	108.5MB/s	00:00
jexec	100%	10KB	21.4MB/s	00:00
orb.idl	100%	640	2.1MB/s	00:00
tr.idl	100%	18KB	18.6MB/s	00:00
javafx-mx.jar	100%	35KB	38.9MB/s	00:00
ant-javafx.jar	100%	1196KB	106.4MB/s	00:00

为了使用方便，我们可以为 JDK 添加环境变量，在 `/usr/profile` 文件末尾追加以下内容，注意 JDK 的环境路径。

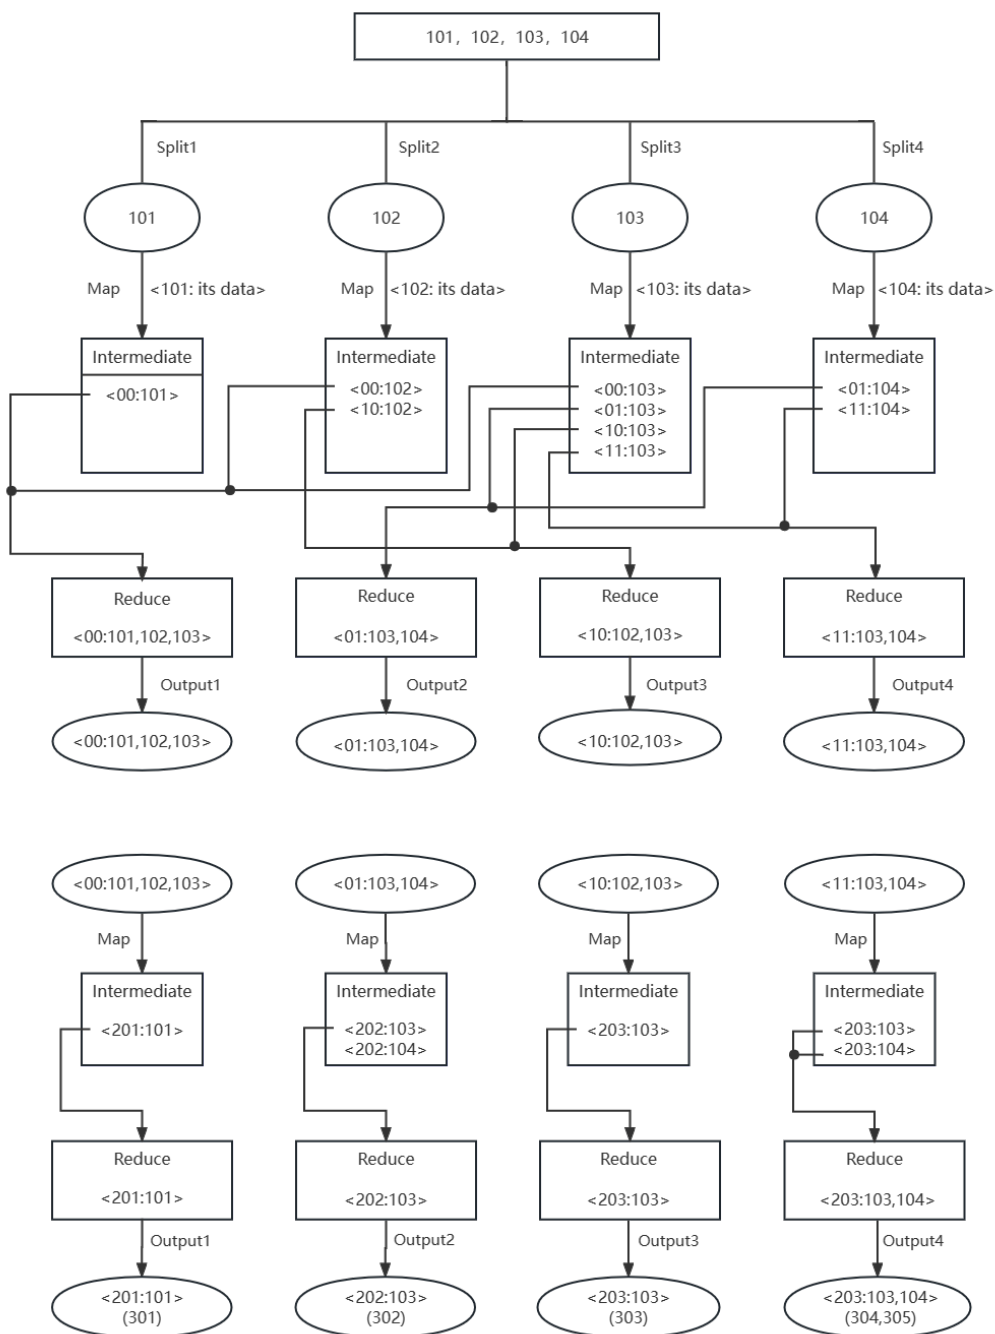
```
for i in /etc/profile.d/*.sh /etc/profile.d/sh.local ; do
    if [ -r "$i" ]; then
        if [ "${-#*i}" != "$-" ]; then
            . "$i"
        else
            . "$i" >/dev/null
        fi
    fi
done

unset i
unset -f pathmunge
#set java environment
export JAVA_HOME=/usr/java/jdk1.8.0_162
export CLASSPATH=.:$JAVA_HOME/lib:$JAVA_HOME/jre/lib:$CLASSPATH
export PATH=$JAVA_HOME/bin:$JAVA_HOME/jre/bin:$PATH
export LC_ALL=zh_CN.utf8
```

编译环境变量后，查询 Java 版本信息，验证配置是否正确：

```
[root@masterNode1 ~]# source /etc/profile
[root@masterNode1 ~]# java -version
java version "1.8.0_162"
Java(TM) SE Runtime Environment (build 1.8.0_162-b12)
Java HotSpot(TM) 64-Bit Server VM (build 25.162-b12, mixed mode)
```

## 第四次作业



上图，主题层多边形从 101 编码到 104 被提交到主机，然后在将所有资源随机分为四项任务后接收多边形 ID 和它们的数据，每个 map 任务产生一个由网格的中间键值对组成的 ID 和一个多边形的 ID，它指示多边形与网格相交。随后，中间键值具有相同网格 ID 的对被输入到相同的 reduce 任务中。这些多边形与网格执行交叉测试，那些与网格不相交的多边形被从键值对列表中删除。类似地，裁剪层的网格索引多边形被生成。网格索引的建立结果如下图。这个 map 任务的输出是剪辑层多边形的 ID 和有可能与前者相交的主题层的 ID 的中间键值对。在 reduce 阶段，每个任务都有具有相同剪辑层多边形 ID 的中间键值对。reduce 任务的输出是裁剪层多边形与主题多边形相交得到的多边形。